

TARUN REDDY GADDAM

📍 Brooklyn, USA ✉ gaddamtaranreddy0@gmail.com ☎ +1(913)899-7287 📄 in/tarunreddy-gaddam-06819b303

SUMMARY

Data Engineer with 2+ years of professional experience designing and operating cloud-native data pipelines, real-time streaming systems, and AI-augmented ETL workflows. Built production pipelines on AWS (Glue, S3, Redshift, EC2), Apache Kafka, and Apache Spark processing millions of records daily. Hands-on expertise integrating HuggingFace transformer models directly into ingestion pipelines — enabling intelligent enrichment at scale. Completed an M.S. in Information Technology Management (Aug 2025, St. Francis College, Brooklyn) while building a portfolio of open-source data engineering projects.

EXPERIENCE

Data Engineer

DEEP IT LLC

June 2025 - Present, Remote

- Architected and deployed 6 end-to-end AWS Glue ETL pipelines ingesting data from Oracle, SQL Server, and flat-file sources into a centralized S3 data lake, reducing data latency from 24 hours to under 90 minutes (62% improvement).
- Designed and optimized 40+ complex SQL queries (CTEs, window functions, partitioned tables) against enterprise Oracle and SQL Server databases supporting 500K+ daily records.
- Built a reusable Python framework for parameterized Glue job execution, cutting new pipeline development time by 35% and eliminating boilerplate across 4 project teams.
- Implemented data quality validation layer using Great Expectations, catching 12,000+ bad records per week before they reached downstream BI dashboards.
- Led migration of 3 legacy on-prem SQL Server ETL jobs to AWS Glue, reducing infrastructure cost by \$18K/year and improving pipeline reliability from 94% to 99.8% uptime.
- Collaborated with data scientists to pipeline HuggingFace NLP models into ingestion workflows, enabling real-time sentiment tagging of 10K+ customer records daily.
- Mentored 2 junior engineers on Spark DataFrame API and AWS Glue best practices; authored internal wiki documentation adopted team-wide.

Junior Data Engineer

Infosys BPM Ltd.

August 2022 - December 2023, Hyderabad, India

- Developed and maintained 15+ Python-based ETL pipelines processing financial transaction data for a US banking client, handling 2M+ records per batch run.
- Optimized slow-running PL/SQL stored procedures and batch jobs, achieving average 40% query performance improvement across critical reporting pipelines.
- Built Apache Kafka producer/consumer pipelines for real-time transaction event streaming, reducing fraud alert latency from 8 minutes to under 30 seconds.
- Containerized data pipeline services using Docker and Docker Compose, enabling consistent environment parity between development, staging, and production.
- Created Power BI dashboards tracking 5 KPIs for operations leadership, consolidating data from 3 source systems into a single governed reporting layer.

PROJECT

TechPulse — Real-Time AI Trend Intelligence Pipeline

github.com/tarunreddy25/TechPulse

- Built end-to-end real-time pipeline ingesting ~120 live records per run from GitHub Trending, Reddit, and Hacker News APIs simultaneously; streamed through Kafka on Docker, fully decoupling ingestion from AI processing.
- Scored every headline using DistilBERT transformer model; discovered 70%+ negative sentiment in tech news discourse and detected real-world events (Java 26 release, OpenAI Sora shutdown) in real time.
- Shipped 4-chart Power BI dashboard revealing AI as dominant keyword by 4× margin.

NYC Taxi Analytics — Large-Scale Data Engineering & Quality

github.com/tarunreddy25/nyc-taxi-project

- Ingested and processed 2.96 million rows of real NYC Yellow Taxi trip data; designed schema, ingestion, and transformation layers end-to-end.
- Identified and removed 240,819 invalid records (including trips > 300,000 miles) — demonstrating production-grade data quality instincts.
- Answered 5 business questions via SQL: peak-hour demand, revenue by weekday, tip behavior segmentation, borough-level trip categorization, and driver utilization rates.

EDUCATION

M.S. Information Technology Management

St. Francis College, Brooklyn, NY • Brooklyn, NY • August 2025

• Completed full-time during Dec 2023 – Aug 2025; built personal data engineering projects and open-source portfolio throughout.

B.Tech Computer Science & Engineering

JNTU Hyderabad, India • Hyderabad, India • August 2022

SKILLS

Python, SQL, Bash, Scala

AWS Glue, S3, EC2, Redshift, Athena, IAM

Apache Kafka, Apache Spark, Apache Airflow, DuckDB, PostgreSQL, SQL Server, Oracle

HuggingFace Transformers, DistilBERT, Sentiment Analysis, NLP Pipelines, MLOps basics

Docker, Docker Compose, Git, GitHub, Great Expectations, JIRA, Confluence

Power BI, Pandas, Matplotlib

Data Lakehouse, Streaming Architecture, Data Modeling, Data Quality, Data Governance
